GLM-5.2: One MoE Decoder Layer

Representative sparse layer with a full DSA indexer (for example, layer 7 of 78, 1-based); batch size fixed at 1
Module paths are relative to model.layers.<layer_idx> unless prefixed with model.

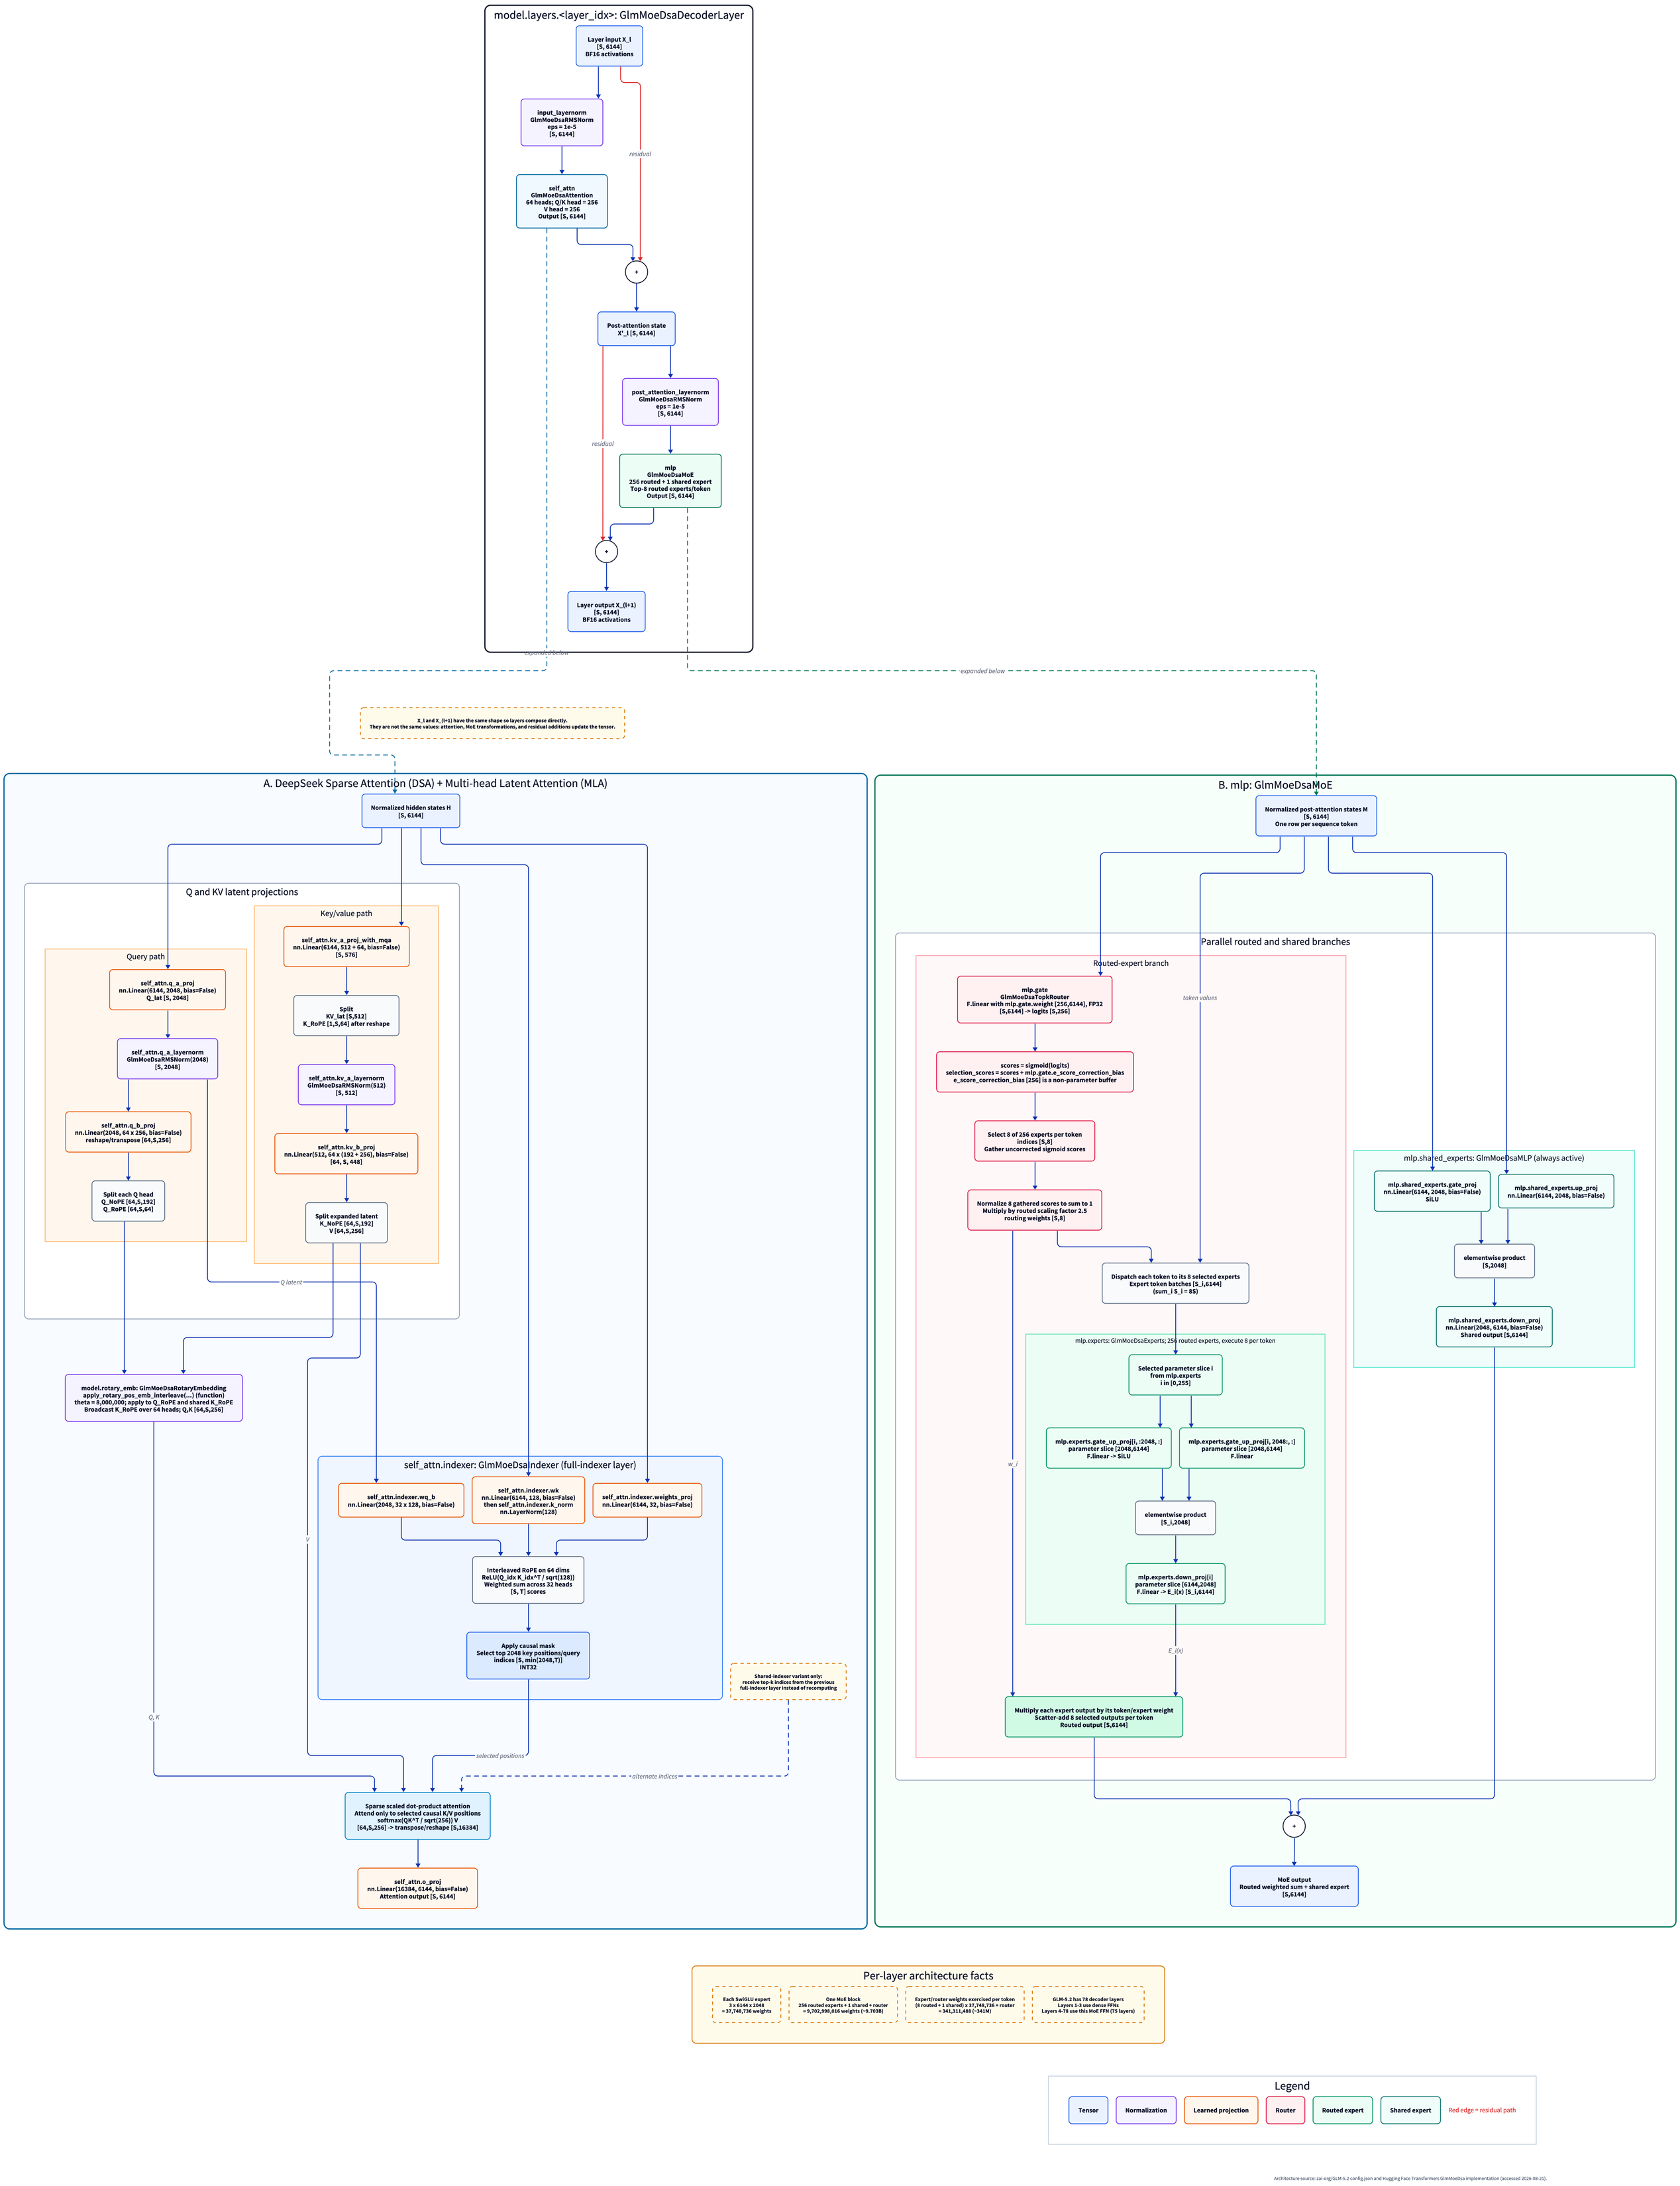

Per-layer architecture facts

Each SwiGLU expert	One MoE block	Expert/router weights exercised per token	GLM-5.2 has 78 decoder layers
3 x 6144 x 2048 = 37,748,736 weights	256 routed experts + 1 shared + router = 9,702,998,016 weights (~9.703B)	(8 routed + 1 shared) x 37,748,736 + router = 341,311,488 (~341M)	Layers 1-3 use dense FFNs Layers 4-78 use this MoE FFN (75 layers)

Legend

Tensor · Normalization · Learned projection · Router · Routed expert · Shared expert · Red edge = residual path

Architecture source: zai-org/GLM-5.2 config.json and Hugging Face Transformers GlmMoeDsa implementation (accessed 2026-08-21).